



# Graphical Models: Introduction

Thomas Hamelryck, April 2015

Bioinformatics center

Department of Biology

University of Copenhagen

# Some history (1)

- Expert systems were developed in the 60s-70s and commercialized in the 80s.
- Expert system=knowledge base+inference engine
- Typical example application is disease diagnosis based on a list of symptoms.
- Expert systems are based on “IF ... THEN ...” rules:
  - Decision trees
  - Rule based or production systems (PROLOG, LISP)

# Some history (1a)

- Example

IF the animal has hair THEN it is a mammal

IF the animal gives milk THEN it is a mammal

IF the animal lays eggs and flies THEN it is a bird



# Some history (2)

- Expert systems use **logical deduction**
  - If A is true, then B is true
    - A is true, therefore, B is true
    - B is false, therefore A is false
- A single unexpected or absent symptom can corrupt disease diagnosis completely.
- Difficult to deal with **unobserved variables** (for example, blood pressure that was not measured).
- The set of rules can become huge.

# Some history (3)

- Expert systems do not use **logical induction**
  - If A is true, then B is true
    - B is true, therefore, A becomes more plausible
    - A is false, therefore, B becomes less plausible
- Rule based systems need to cope with **degrees of uncertainty**.
  - The AI community tried various solutions. Each rule is associated with a **certainty factor (CF)**. These factors are combined according to a certain **algebra**.
    - fuzzy logic, belief functions,...

# Some history (3a)



Pierre-Simon Laplace  
(1749-1827)

## ■ Example

IF headache, fever THEN influenza (CF=0.7)

IF influenza THEN sneezing (CF=0.9)

IF influenza THEN weakness (CF=0.6)

## ■ CF algebra

■ What is  $CF(\text{influenza} \mid \text{sneezing, no headache})$ ?

■ These *ad hoc* algebra's were proven to be **inconsistent**

- Only Bayesian probability will do if we follow some elementary desiderata (Richard T. Cox, 1946,1961; Edwin T. Jaynes, 2003, see Bishop, p. 21)

# Some history (4)

- Despite this, using probabilities as certainty factor was seen as problematic, mainly because
  - According to the **frequentist interpretation**, probabilities are essentially frequencies in a large number of trials. Often this interpretation is meaningless in expert systems.
  - Full probability distributions over many variables are **intractable**.
- Then came Judea Pearl's "Probabilistic inference in intelligent systems", 1988
  - **Bayesian** interpretation of probability
  - Efficient local computations on **graphs**

# Background references

## ■ Probability theory

- *Probability theory – the logic of science.* (2003) Edwin T. Jaynes (Cambridge university press)
- *The algebra of probable inference.* (1961) Richard T. Cox (Johns Hopkins univ. press)

## ■ Expert systems & BNs

- *Probabilistic inference in intelligent systems.* (1988) Judea Pearl (Morgan Kauffman)
- *Probabilistic networks and expert systems.* (1999) Cowell, Dawid, Lauritzen, Spiegelhalter (Springer)
- *Bayesian networks and decision graphs.* (2007) Finn V. Jensen, Thomas D. Nielsen (Springer)



# Some preliminaries (1)

- Conditional probability tables (CPTs)
  - Discrete random variables with finite number of states
  - Probability of one variable conditional on one or more variables
    - Example:  $P(\text{car ownership}|\text{size of income})$

	No car	Second hand car	New car
Low income	0.2	0.4	0.4
High income	0.1	0.3	0.6

# Some preliminaries (2)

- Discrete random variables with finite range
  - K possible states
  - 1-of-K representation
    - $\mathbf{x}$  is a K-dimensional vector of binary indicators
    - Example
      - $\mathbf{x}=(0,0,1,0)$  indicates the third state (out of  $K=4$ )
  - If the probabilities of the K states are  $(\mu_1, \mu_2, \dots, \mu_K)$  then:

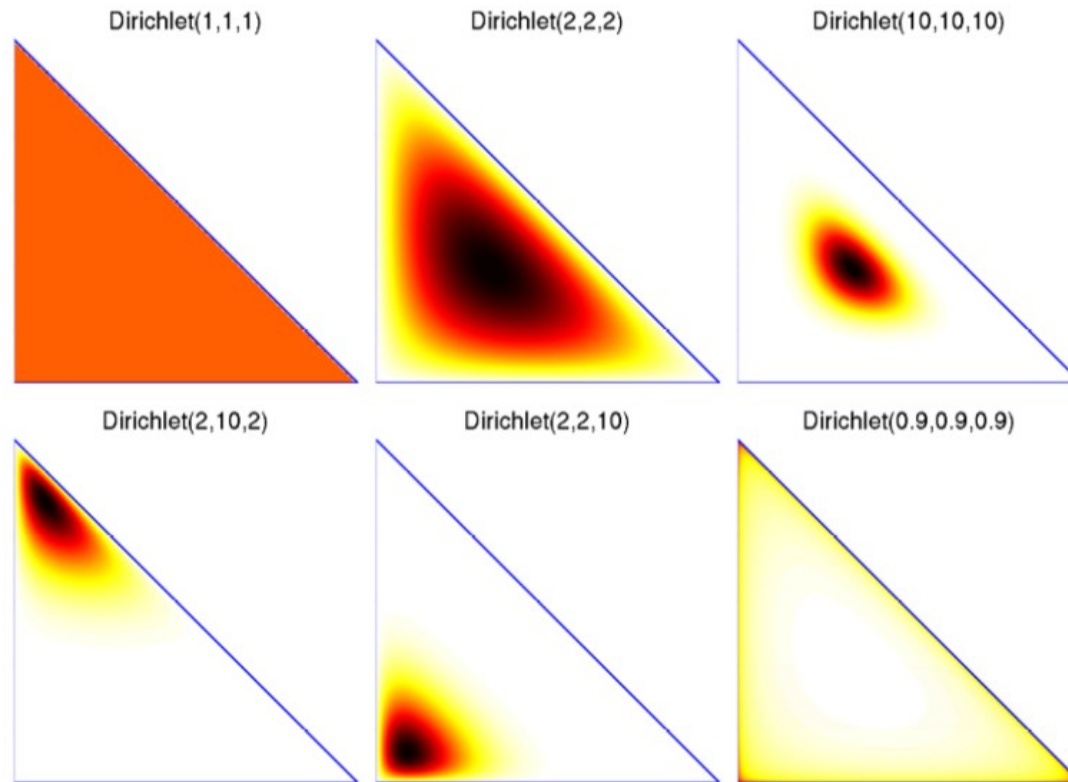
$$P(\mathbf{x}|\boldsymbol{\mu}) = \prod_{k=1}^K \mu_k^{x_k}$$

# Some preliminaries (3a)

- Dirichlet distribution  $P(\mu_1, \dots, \mu_K | \alpha_1, \dots, \alpha_K) = \frac{1}{C(\boldsymbol{\alpha})} \prod_{k=1}^K \mu_k^{\alpha_k - 1}$ 
  - Parameter  $\boldsymbol{\alpha}$ 
    - K-dimensional vector of reals > 0
  - Probability distribution over vectors  $\boldsymbol{\mu}$ 
    - K-dimensional
    - Positive components that sum to one
    - Probability vector
  - Example of a probability distribution on a “special” manifold
    - K-dimensional simplex
      - Generalization of a triangle



# Some preliminaries (3b)

- Dirichlet distributions with  $K=3$



# Some preliminaries (4)

- Distributions on the sphere ( $S^2$ ) and the torus ( $T^2$ )
  - Probabilistic models of protein structure



---

# **PATTERN RECOGNITION AND MACHINE LEARNING**


## **CHAPTER 8: GRAPHICAL MODELS**

---

# Bayesian Networks

---

Directed Acyclic Graph (DAG)



$$p(a, b, c) = p(c|a, b)p(a, b) = p(c|a, b)p(b|a)p(a)$$


Product rule:  $p(x_1, \dots, x_K) = p(x_K|x_1, \dots, x_{K-1}) \dots p(x_2|x_1)p(x_1)$

---

# Bayesian Networks

---

$$p(x_1, \dots, x_7) = p(x_1)p(x_2)p(x_3)p(x_4|x_1, x_2, x_3) \\ p(x_5|x_1, x_3)p(x_6|x_4)p(x_7|x_4, x_5)$$




General Factorization

$$p(\mathbf{x}) = \prod_{k=1}^K p(x_k | \text{pa}_k)$$



# Bayesian Curve Fitting (1)

---



Polynomial


$$y(x, \mathbf{w}) = \sum_{j=0}^M w_j x^j$$

$$p(\mathbf{t}, \mathbf{w}) = p(\mathbf{w}) \prod_{n=1}^N p(t_n | y(\mathbf{w}, x_n))$$

# Bayesian Curve Fitting (2)

---

$$p(\mathbf{t}, \mathbf{w}) = p(\mathbf{w}) \prod_{n=1}^N p(t_n | y(\mathbf{w}, x_n))$$




# Bayesian Curve Fitting (3)

---

Input variables and explicit hyperparameters

$$p(\mathbf{t}, \mathbf{w} | \mathbf{x}, \alpha, \sigma^2) = p(\mathbf{w} | \alpha) \prod_{n=1}^N p(t_n | \mathbf{w}, x_n, \sigma^2).$$




# Bayesian Curve Fitting—Learning

---

Condition on data

$$p(\mathbf{w}|\mathbf{t}) \propto p(\mathbf{w}) \prod_{n=1}^N p(t_n|\mathbf{w}) \quad (\text{Bayes})$$

Shaded=observed  
Not shaded=hidden, latent




# Bayesian Curve Fitting—Prediction

(product rule)

Predictive distribution:  $p(\hat{t}|\hat{x}, \mathbf{x}, \mathbf{t}, \alpha, \sigma^2) \propto \int p(\hat{t}, \mathbf{t}, \mathbf{w}|\hat{x}, \mathbf{x}, \alpha, \sigma^2) d\mathbf{w}$

where


$$p(\hat{t}, \mathbf{t}, \mathbf{w}|\hat{x}, \mathbf{x}, \alpha, \sigma^2) = \left[ \prod_{n=1}^N p(t_n|x_n, \mathbf{w}, \sigma^2) \right] p(\mathbf{w}|\alpha)p(\hat{t}|\hat{x}, \mathbf{w}, \sigma^2)$$



# Generative Models

---

Causal process for generating images




Ancestral sampling:  
p. 365, Bishop

---

# Discrete Variables (1)

---

General joint distribution:  $K^2 - 1$  parameters



$$p(\mathbf{x}_1, \mathbf{x}_2 | \boldsymbol{\mu}) = \prod_{k=1}^K \prod_{l=1}^K \mu_{kl}^{x_{1k} x_{2l}}$$

Independent joint distribution:  $2(K - 1)$  parameters



$$\hat{p}(\mathbf{x}_1, \mathbf{x}_2 | \boldsymbol{\mu}) = \prod_{k=1}^K \mu_{1k}^{x_{1k}} \prod_{l=1}^K \mu_{2l}^{x_{2l}}$$

Independence assumptions lead to fewer parameters

---

# Discrete Variables (2)

---

General joint distribution over  $M$  variables:

$K^M - 1$  parameters


$M$ -node Markov chain:  $K - 1 + (M - 1)K(K - 1)$   
parameters





# Discrete Variables: Bayesian Parameters (1)

---




$$p(\{\mathbf{x}_m, \mu_m\}) = p(\mathbf{x}_1 | \mu_1) p(\mu_1) \prod_{m=2}^M p(\mathbf{x}_m | \mathbf{x}_{m-1}, \mu_m) p(\mu_m)$$

$$p(\mu_m) = \text{Dir}(\mu_m | \alpha_m)$$

---

## Discrete Variables: Bayesian Parameters (2)


---



$$p(\{x_m\}, \mu_1, \mu) = p(x_1 | \mu_1) p(\mu_1) \prod_{m=2}^M p(x_m | x_{m-1}, \mu) p(\mu)$$

# Parameterized Conditional Distributions

---



If  $x_1, \dots, x_M$  are discrete,  $K$ -state variables,  $p(y = 1|x_1, \dots, x_M)$  in general has  $O(K^M)$  parameters.

The parameterized form

$$p(y = 1|x_1, \dots, x_M) = \sigma \left( w_0 + \sum_{i=1}^M w_i x_i \right) = \sigma(\mathbf{w}^T \mathbf{x})$$

requires only  $M + 1$  parameters

logistic sigmoid =  $1/(1 + \exp(-a))$


# Linear-Gaussian Models

---

## Directed Graph

$$p(x_i | \text{pa}_i) = \mathcal{N} \left( x_i \left| \sum_{j \in \text{pa}_i} w_{ij} x_j + b_i, v_i \right. \right)$$

Mean      Standard deviation




Each node is Gaussian, the mean is a linear function of the parents.

## Vector-valued Gaussian Nodes

$$p(\mathbf{x}_i | \text{pa}_i) = \mathcal{N} \left( \mathbf{x}_i \left| \sum_{j \in \text{pa}_i} \mathbf{W}_{ij} \mathbf{x}_j + \mathbf{b}_i, \Sigma_i \right. \right)$$

Covariance Matrix



# Conditional Independence

---

$a$  is independent of  $b$  given  $c$

$$p(a|b, c) = p(a|c)$$

Equivalently

$$\begin{aligned} p(a, b|c) &= p(a|b, c)p(b|c) \\ &= p(a|c)p(b|c) \end{aligned}$$

Notation


$$a \perp\!\!\!\perp b \mid c$$

Bayesian networks encode conditional independences

---

# Conditional Independence: Example 1

---



Path is "unblocked".

$$p(a, b, c) = p(a|c)p(b|c)p(c)$$


$$p(a, b) = \sum_c p(a|c)p(b|c)p(c)$$

$$a \not\perp b \mid \emptyset$$

# Conditional Independence: Example 1

---

c is tail-to-tail




Observing c "blocks" the path  
and makes a and b independent

$$\begin{aligned} p(a, b|c) &= \frac{p(a, b, c)}{p(c)} \\ &= p(a|c)p(b|c) \end{aligned}$$

$$a \perp\!\!\!\perp b \mid c$$

# Conditional Independence: Example 2

---



$$p(a, b, c) = p(a)p(c|a)p(b|c)$$

$$p(a, b) = p(a) \sum_c p(c|a)p(b|c) = p(a)p(b|a)$$

$$a \not\perp b \mid \emptyset$$


Path is "unblocked".

---



# Conditional Independence: Example 2

---



$$\begin{aligned} p(a, b|c) &= \frac{p(a, b, c)}{p(c)} \\ &= \frac{p(a)p(c|a)p(b|c)}{p(c)} \\ &= p(a|c)p(b|c) \end{aligned}$$


Observing  $c$  "blocks" the path  
and makes  $a$  and  $b$  independent

$$a \perp\!\!\!\perp b \mid c$$

---

# Conditional Independence: Example 3

---



$$p(a, b, c) = p(a)p(b)p(c|a, b)$$

$$p(a, b) = p(a)p(b)$$

$$a \perp\!\!\!\perp b \mid \emptyset$$


Path is "blocked".

Note: this is the opposite of Example 1, with  $c$  unobserved.

---

# Conditional Independence: Example 3

---



$c$  is head-to-head

$$\begin{aligned} p(a, b|c) &= \frac{p(a, b, c)}{p(c)} \\ &= \frac{p(a)p(b)p(c|a, b)}{p(c)} \end{aligned}$$

$$a \not\perp b \mid c$$

Observing  $c$  "unblocks" the path  
and makes  $a$  and  $b$  dependent

Note: this is the opposite of Example 1, with  $c$  observed.

---

# “Am I out of fuel?”

---


A very lousy fuel gauge:

$$p(G = 1 | B = 1, F = 1) = 0.8$$

$$p(G = 1 | B = 1, F = 0) = 0.2$$

$$p(G = 1 | B = 0, F = 1) = 0.2$$

$$p(G = 1 | B = 0, F = 0) = 0.1$$



Priors:

$$p(B = 1) = 0.9$$

$$p(F = 1) = 0.9$$

and hence

$$p(F = 0) = 0.1$$

$B$  = Battery (0=flat, 1=fully charged)


$F$  = Fuel Tank (0=empty, 1=full)

$G$  = Fuel Gauge Reading  
(0=empty, 1=full)

---

# “Am I out of fuel?”

---



Probability of empty tank given gauge says its empty:


$$\begin{aligned} p(F = 0|G = 0) &= \frac{p(G = 0|F = 0)p(F = 0)}{p(G = 0)} \\ &\simeq 0.257 \end{aligned}$$

Probability of an empty tank increased by observing  $G = 0$ .

---

# “Am I out of fuel?”

---



Now also given that the battery is empty:

$$\begin{aligned} p(F = 0 | G = 0, B = 0) &= \frac{p(G = 0 | B = 0, F = 0)p(F = 0)}{\sum_{F \in \{0,1\}} p(G = 0 | B = 0, F)p(F)} \\ &\simeq 0.111 \end{aligned}$$

Decreased!

Probability of an empty tank reduced by observing  $B = 0$ .

This referred to as “explaining away” of  $G$  by  $B$ .

---

Fuel and battery are not independent because the gauge is observed.


# D-separation

---


- $A$ ,  $B$ , and  $C$  are non-intersecting subsets of nodes in a directed graph.
  - A path from  $A$  to  $B$  is **blocked** if it contains a node such that either
    - a) the arrows on the path meet either **head-to-tail** or **tail-to-tail** at the node, and the node is in the set  $C$ , or
    - b) the arrows meet **head-to-head** at the node, and neither the node, nor any of its descendants, are in the set  $C$ .
  - If all paths from  $A$  to  $B$  are blocked,  $A$  is said to be **d-separated** from  $B$  by  $C$ .
  - If  $A$  is d-separated from  $B$  by  $C$ , the joint distribution over all variables in the graph satisfies  $A \perp\!\!\!\perp B \mid C$ .
-

# D-separation: Example

---



$$a \not\perp b \mid c$$



$$a \perp b \mid f$$

What about  $\mid e$ ?

What about  $\mid$  empty set?


---



# D-separation: I.I.D. Data

---

Independent identically distributed

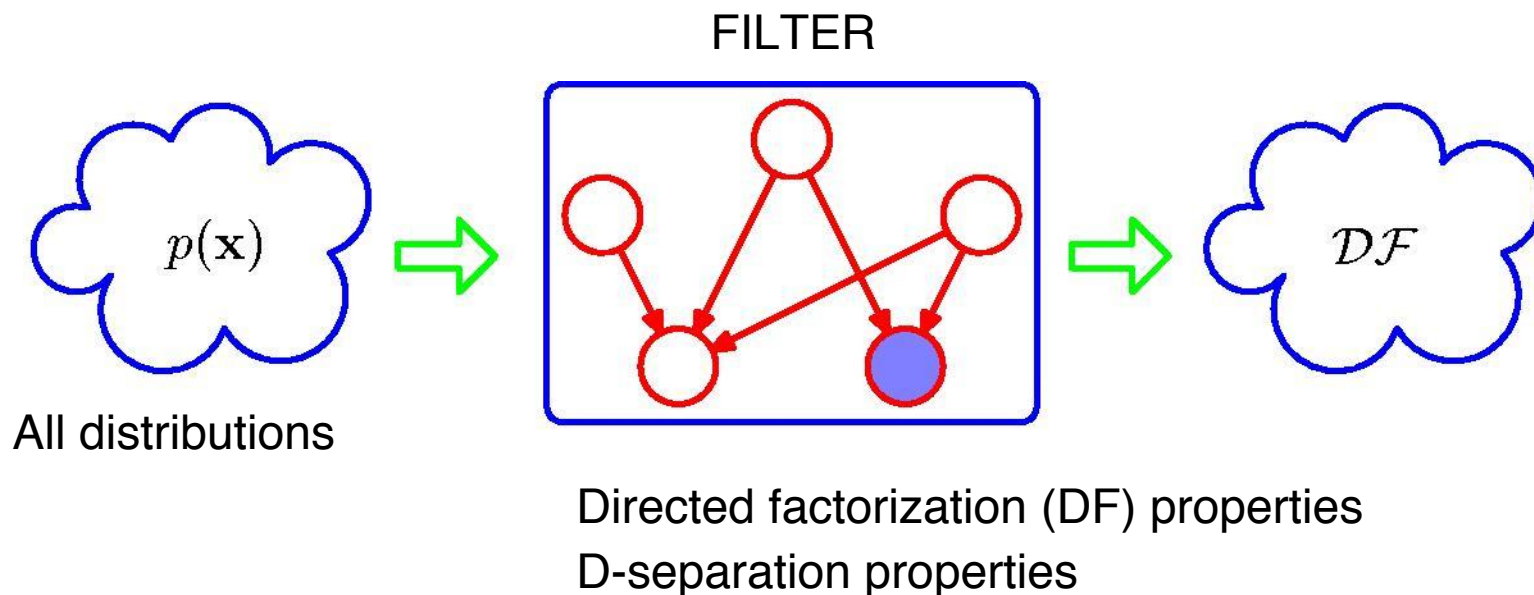


$$p(\mathcal{D}|\mu) = \prod_{n=1}^N p(x_n|\mu)$$

d-separated by  $\mu$


$$p(\mathcal{D}) = \int_{-\infty}^{\infty} p(\mathcal{D}|\mu)p(\mu) d\mu \neq \prod_{n=1}^N p(x_n) \quad \text{not d-separated by } \emptyset$$

# Directed Graphs as Distribution Filters



# The Markov Blanket

---




Dependent only on:  
Parents  
Children  
Parents of children

$$\begin{aligned} p(\mathbf{x}_i | \mathbf{x}_{\{j \neq i\}}) &= \frac{p(\mathbf{x}_1, \dots, \mathbf{x}_M)}{\int p(\mathbf{x}_1, \dots, \mathbf{x}_M) d\mathbf{x}_i} \\ &= \frac{\prod_k p(\mathbf{x}_k | \text{pa}_k)}{\int \prod_k p(\mathbf{x}_k | \text{pa}_k) d\mathbf{x}_i} \end{aligned}$$

Factors independent of  $\mathbf{x}_i$  cancel  
between numerator and denominator.

# Converting Directed to Undirected Graphs (1)


---



$$p(\mathbf{x}) = \underbrace{p(x_1)p(x_2|x_1)}_{\text{red bracket}} p(x_3|x_2) \cdots p(x_N|x_{N-1})$$

$$p(\mathbf{x}) = \frac{1}{Z} \psi_{1,2}(x_1, x_2) \psi_{2,3}(x_2, x_3) \cdots \psi_{N-1,N}(x_{N-1}, x_N)$$


Three red double-headed arrows connect the underbraced term in the equation above to the first three  $\psi$  terms in this equation: from the bracket to  $\psi_{1,2}$ , from  $p(x_3|x_2)$  to  $\psi_{2,3}$ , and from  $p(x_N|x_{N-1})$  to  $\psi_{N-1,N}$ .



# Inference in Graphical Models

---

We observe  $y$ . What is  $p(x|y)$ ?



$$p(y) = \sum_{x'} p(y|x')p(x')$$

Sum rule


$$p(x|y) = \frac{p(y|x)p(x)}{p(y)}$$

Bayes' formula

---

# Inference on a Chain

---



$$p(\mathbf{x}) = \frac{1}{Z} \psi_{1,2}(x_1, x_2) \psi_{2,3}(x_2, x_3) \cdots \psi_{N-1,N}(x_{N-1}, x_N)$$

$$p(x_n) = \sum_{x_1} \cdots \sum_{x_{n-1}} \sum_{x_{n+1}} \cdots \sum_{x_N} p(\mathbf{x})$$

Naive summation:


For  $N$  variables with  $K$  states:  $K^{N-1}$  terms

Last term is now  
only function of  
 $x_{N-1}$

---

Perform summation first over  $x_N$ , and save operations:  $ab+ac=a(b+c)$

# Inference on a Chain




$$p(x_n) = \frac{1}{Z} \underbrace{\left[ \sum_{x_{n-1}} \psi_{n-1,n}(x_{n-1}, x_n) \cdots \left[ \sum_{x_1} \psi_{1,2}(x_1, x_2) \right] \cdots \right]}_{\mu_\alpha(x_n)} \underbrace{\left[ \sum_{x_{n+1}} \psi_{n,n+1}(x_n, x_{n+1}) \cdots \left[ \sum_{x_N} \psi_{N-1,N}(x_{N-1}, x_N) \right] \cdots \right]}_{\mu_\beta(x_n)}$$

This trick can be seen as passing local messages around in the graph.

# Inference on a Chain

---



$$\mu_\alpha(x_n) = \sum_{x_{n-1}} \psi_{n-1,n}(x_{n-1}, x_n) \left[ \sum_{x_{n-2}} \cdots \right]$$

$$= \sum_{x_{n-1}} \psi_{n-1,n}(x_{n-1}, x_n) \mu_\alpha(x_{n-1}).$$

$$\mu_\beta(x_n) = \sum_{x_{n+1}} \psi_{n,n+1}(x_n, x_{n+1}) \left[ \sum_{x_{n+2}} \cdots \right]$$

$$= \sum_{x_{n+1}} \psi_{n,n+1}(x_n, x_{n+1}) \mu_\beta(x_{n+1}).$$


---

These messages can be evaluated recursively.



# Inference on a Chain

---



$$\mu_\alpha(x_2) = \sum_{x_1} \psi_{1,2}(x_1, x_2) \qquad \mu_\beta(x_{N-1}) = \sum_{x_N} \psi_{N-1,N}(x_{N-1}, x_N)$$

$$Z = \sum_{x_n} \mu_\alpha(x_n) \mu_\beta(x_n)$$

Start of the recursion, and calculation of the normalization constant-.

---

# Inference on a Chain

---

To compute local marginals:

- Compute and store all forward messages,  $\mu_\alpha(x_n)$ .
- Compute and store all backward messages,  $\mu_\beta(x_n)$ .
- Compute  $Z$  at any node  $x_m$
- Compute

$$p(x_n) = \frac{1}{Z} \mu_\alpha(x_n) \mu_\beta(x_n)$$


for all variables required.

---


# Trees

---


Undirected Tree



Directed Tree




Polytree



# Factor Graphs

---




$$p(\mathbf{x}) = f_a(x_1, x_2) f_b(x_1, x_2) f_c(x_2, x_3) f_d(x_3)$$

$$p(\mathbf{x}) = \prod_s f_s(\mathbf{x}_s)$$


---

# Factor Graphs from Directed Graphs


---



$$p(\mathbf{x}) = p(x_1)p(x_2) \\ p(x_3|x_1, x_2)$$



$$f(x_1, x_2, x_3) = \\ p(x_1)p(x_2)p(x_3|x_1, x_2)$$



$$f_a(x_1) = p(x_1)$$

$$f_b(x_2) = p(x_2)$$

$$f_c(x_1, x_2, x_3) = p(x_3|x_1, x_2)$$


---

# The Sum-Product Algorithm (1)

---

## Objective:

- i. to obtain an efficient, exact inference algorithm for finding marginals;
- ii. in situations where several marginals are required, to allow computations to be shared efficiently.


## Key idea: Distributive Law

$$ab + ac = a(b + c)$$

---

# The Sum-Product Algorithm (2)

---




$\text{ne}(x)$ =factor nodes attached to  $x$

$X_s$ =variable nodes attached to  $x$  via factor  $s$

$$p(x) = \sum_{\mathbf{x} \setminus x} p(\mathbf{x})$$

$$p(\mathbf{x}) = \prod_{s \in \text{ne}(x)} F_s(x, X_s)$$

# The Sum-Product Algorithm (3)



(from the previous two formula's,  
and interchanging sum and product)

$$\begin{aligned}
 p(x) &= \prod_{s \in \text{ne}(x)} \left[ \sum_{X_s} F_s(x, X_s) \right] \\
 &= \prod_{s \in \text{ne}(x)} \mu_{f_s \rightarrow x}(x).
 \end{aligned}$$


$$\mu_{f_s \rightarrow x}(x) \equiv \sum_{X_s} F_s(x, X_s)$$

Again, this can be viewed as messages, from factor nodes  $f_s$  to variable node  $x$ .



# The Sum-Product Algorithm (4)


---



How is  $F_s$  calculated?

$$F_s(x, X_s) = f_s(x, x_1, \dots, x_M) G_1(x_1, X_{s1}) \dots G_M(x_M, X_{sM})$$


# The Sum-Product Algorithm (5)



$$\begin{aligned} \mu_{f_s \rightarrow x}(x) &= \sum_{x_1} \dots \sum_{x_M} f_s(x, x_1, \dots, x_M) \prod_{m \in \text{ne}(f_s) \setminus x} \left[ \sum_{X_{sm}} G_m(x_m, X_{sm}) \right] \\ &= \sum_{x_1} \dots \sum_{x_M} f_s(x, x_1, \dots, x_M) \prod_{m \in \text{ne}(f_s) \setminus x} \mu_{x_m \rightarrow f_s}(x_m) \end{aligned}$$

Second type of message! This time from variable to factor node!

# The Sum-Product Algorithm (6)




$$\begin{aligned}
 \mu_{x_m \rightarrow f_s}(x_m) &\equiv \sum_{X_{sm}} G_m(x_m, X_{sm}) = \sum_{X_{sm}} \prod_{l \in \text{ne}(x_m) \setminus f_s} F_l(x_m, X_{ml}) \\
 &= \prod_{l \in \text{ne}(x_m) \setminus f_s} \mu_{f_l \rightarrow x_m}(x_m)
 \end{aligned}$$

Messages from variable nodes to factor nodes.


# The Sum-Product Algorithm (7)

---

## Initialization



Messages from  
variable to factor node



Messages from  
factor to variable node

---

# The Sum-Product Algorithm (8)


---

To compute local marginals:

- Pick an arbitrary node as root
  - Compute and propagate messages from the leaf nodes to the root, storing received messages at every node.
  - Compute and propagate messages from the root to the leaf nodes, storing received messages at every node.
  - Compute the product of received messages at each node for which the marginal is required, and normalize if necessary.
-

# Sum-Product: Example (1)

---




$$\tilde{p}(\mathbf{x}) = f_a(x_1, x_2) f_b(x_2, x_3) f_c(x_2, x_4)$$

---

# Sum-Product: Example (2)

---



$$\mu_{x_1 \rightarrow f_a}(x_1) = 1$$

$$\mu_{f_a \rightarrow x_2}(x_2) = \sum_{x_1} f_a(x_1, x_2)$$

$$\mu_{x_4 \rightarrow f_c}(x_4) = 1$$


$$\mu_{f_c \rightarrow x_2}(x_2) = \sum_{x_4} f_c(x_2, x_4)$$

$$\mu_{x_2 \rightarrow f_b}(x_2) = \mu_{f_a \rightarrow x_2}(x_2) \mu_{f_c \rightarrow x_2}(x_2)$$

$$\mu_{f_b \rightarrow x_3}(x_3) = \sum_{x_2} f_b(x_2, x_3) \mu_{x_2 \rightarrow f_b}(x_2)$$

# Sum-Product: Example (3)

---



$$\mu_{x_3 \rightarrow f_b}(x_3) = 1$$

$$\mu_{f_b \rightarrow x_2}(x_2) = \sum_{x_3} f_b(x_2, x_3)$$

$$\mu_{x_2 \rightarrow f_a}(x_2) = \mu_{f_b \rightarrow x_2}(x_2) \mu_{f_c \rightarrow x_2}(x_2)$$

$$\mu_{f_a \rightarrow x_1}(x_1) = \sum_{x_2} f_a(x_1, x_2) \mu_{x_2 \rightarrow f_a}(x_2)$$

$$\mu_{x_2 \rightarrow f_c}(x_2) = \mu_{f_a \rightarrow x_2}(x_2) \mu_{f_b \rightarrow x_2}(x_2)$$


$$\mu_{f_c \rightarrow x_4}(x_4) = \sum_{x_2} f_c(x_2, x_4) \mu_{x_2 \rightarrow f_c}(x_2)$$

---



# Sum-Product: Example (4)

---



$$\begin{aligned}\tilde{p}(x_2) &= \mu_{f_a \rightarrow x_2}(x_2) \mu_{f_b \rightarrow x_2}(x_2) \mu_{f_c \rightarrow x_2}(x_2) \\ &= \left[ \sum_{x_1} f_a(x_1, x_2) \right] \left[ \sum_{x_3} f_b(x_2, x_3) \right] \\ &\quad \left[ \sum_{x_4} f_c(x_2, x_4) \right] \\ &= \sum_{x_1} \sum_{x_3} \sum_{x_4} f_a(x_1, x_2) f_b(x_2, x_3) f_c(x_2, x_4) \\ &= \sum_{x_1} \sum_{x_3} \sum_{x_4} \tilde{p}(\mathbf{x})\end{aligned}$$

---

# The Max-Sum Algorithm (1)

---

Objective: an efficient algorithm for finding

- i. the value  $\mathbf{x}^{\max}$  that maximises  $p(\mathbf{x})$ ;
- ii. the value of  $p(\mathbf{x}^{\max})$ .

In general, maximum marginals  $\neq$  joint maximum.

	$x = 0$	$x = 1$
$y = 0$	0.3	0.4
$y = 1$	0.3	0.0


$$\arg \max_x p(x, y) = 1 \qquad \arg \max_x p(x) = 0$$

---

# The Max-Sum Algorithm (2)

---

Maximizing over a chain (max-product)



$$\begin{aligned} p(\mathbf{x}^{\max}) &= \max_{\mathbf{x}} p(\mathbf{x}) = \max_{x_1} \dots \max_{x_N} p(\mathbf{x}) \\ &= \frac{1}{Z} \max_{x_1} \dots \max_{x_N} [\psi_{1,2}(x_1, x_2) \dots \psi_{N-1,N}(x_{N-1}, x_N)] \\ &= \frac{1}{Z} \max_{x_1} \left[ \max_{x_2} \left[ \psi_{1,2}(x_1, x_2) \left[ \dots \max_{x_N} \psi_{N-1,N}(x_{N-1}, x_N) \right] \dots \right] \right] \end{aligned}$$

# The Max-Sum Algorithm (3)

---

Generalizes to tree-structured factor graph

$$\max_{\mathbf{x}} p(\mathbf{x}) = \max_{x_n} \prod_{f_s \in \text{ne}(x_n)} \max_{X_s} f_s(x_n, X_s)$$

maximizing as close to the leaf nodes as possible

---

# The Max-Sum Algorithm (4)

---

Max-Product  $\rightarrow$  Max-Sum

For numerical reasons, use

$$\ln \left( \max_{\mathbf{x}} p(\mathbf{x}) \right) = \max_{\mathbf{x}} \ln p(\mathbf{x}).$$

Again, use distributive law

$$\max(a + b, a + c) = a + \max(b, c).$$

---

# The Max-Sum Algorithm (5)

---

## Initialization (leaf nodes)

$$\mu_{x \rightarrow f}(x) = 0 \qquad \mu_{f \rightarrow x}(x) = \ln f(x)$$

## Recursion

$$\mu_{f \rightarrow x}(x) = \max_{x_1, \dots, x_M} \left[ \ln f(x, x_1, \dots, x_M) + \sum_{m \in \text{ne}(f_s) \setminus x} \mu_{x_m \rightarrow f}(x_m) \right]$$

$$\phi(x) = \arg \max_{x_1, \dots, x_M} \left[ \ln f(x, x_1, \dots, x_M) + \sum_{m \in \text{ne}(f_s) \setminus x} \mu_{x_m \rightarrow f}(x_m) \right]$$

$$\mu_{x \rightarrow f}(x) = \sum_{l \in \text{ne}(x) \setminus f} \mu_{f_l \rightarrow x}(x)$$

---

# The Max-Sum Algorithm (6)

---

Termination (root node)

$$p^{\max} = \max_x \left[ \sum_{s \in \text{ne}(x)} \mu_{f_s \rightarrow x}(x) \right]$$
$$x^{\max} = \arg \max_x \left[ \sum_{s \in \text{ne}(x)} \mu_{f_s \rightarrow x}(x) \right]$$

Back-track, for all nodes  $i$  with  $l$  factor nodes to the root ( $l=0$ )


$$\mathbf{x}_l^{\max} = \phi(x_{i,l-1}^{\max})$$

---

# The Max-Sum Algorithm (7)

---

Example: Markov chain





# The Junction Tree Algorithm

---

- *Exact* inference on general graphs.
  - Works by turning the initial graph into a *junction tree* and then running a sum-product-like algorithm.
  - *Intractable* on graphs with large cliques.
-

# Loopy Belief Propagation

---

- Sum-Product on general graphs.
  - Initial unit messages passed across all links, after which messages are passed around until convergence (not guaranteed!).
  - *Approximate* but *tractable* for large graphs.
  - Sometime works well, sometimes not at all.
-

# Sequential Data and Markov Models

Sargur N. Srihari

[srihari@cedar.buffalo.edu](mailto:srihari@cedar.buffalo.edu)

Machine Learning Course:

<http://www.cedar.buffalo.edu/~srihari/CSE574/index.html>

# Sequential Data Examples

- Often arise through measurement of time series
  - Snowfall measurements on successive days
  - Rainfall measurements on successive days
  - Daily values of currency exchange rate
  - Acoustic features at successive time frames in speech recognition
  - Nucleotide base pairs in a strand of DNA
  - Sequence of characters in an English sentence
  - Parts of speech of successive words

# Markov Model – Weather

- The weather of a day is observed as being one of the following:


- State 1: Rainy
- State 2: Cloudy
- State 3: Sunny



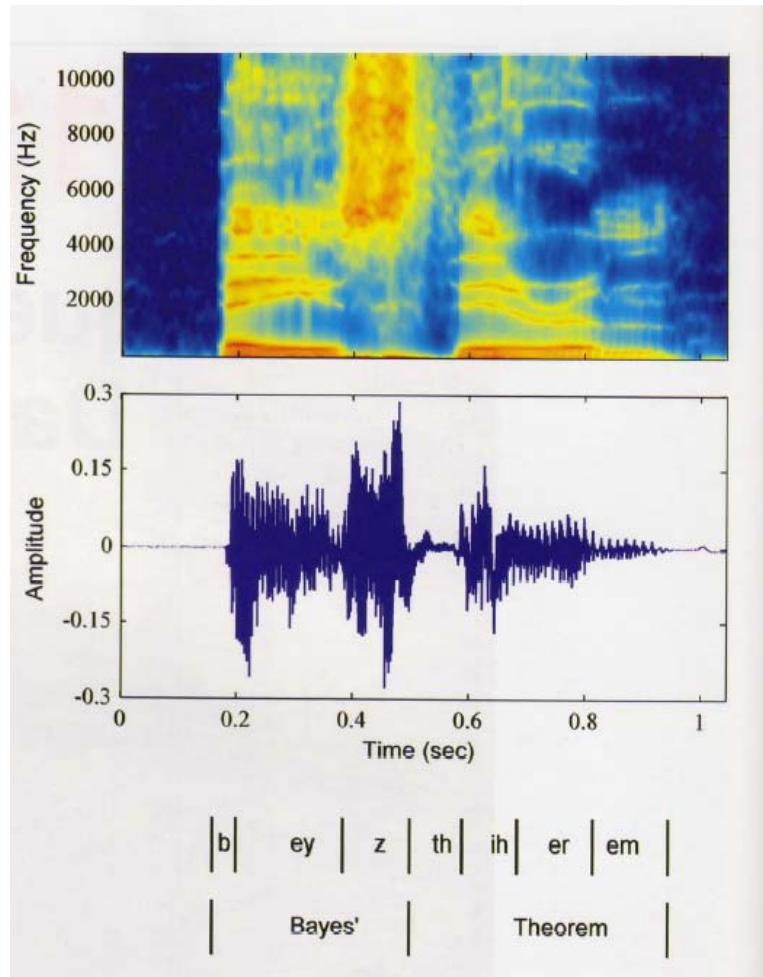
- Weather pattern of a location

		Tomorrow		
		Rain	Cloudy	Sunny
Today	Rain	0.3	0.3	0.4
	Cloudy	0.2	0.6	0.2
	Sunny	0.1	0.1	0.8

# Markov Model – Weather State Diagram



# Sound Spectrogram of Speech




- “Bayes Theorem”
- Plot of the intensity of the spectral coefficients versus time index
- Successive observations of speech spectrum highly correlated (Markov dependency)

# Markov model for the production of spoken words

- States represent phonemes
- Production of word: “cat”
- Represented by states  
/k/ /a/ /t/
- Transitions from
  - /k/ to /a/
  - /a/ to /t/
  - /t/ to a silent state
- Although only correct cat sound is represented by model, perhaps other transitions can be introduced,
  - eg, /k/ followed by /t/

Markov Model  
for word “cat”





# Stationary vs Non-stationary

- Stationary: Data evolves over time but distribution remains same
  - e.g., dependence of current word over previous word remains constant
- Non-stationary: Generative distribution itself changes over time

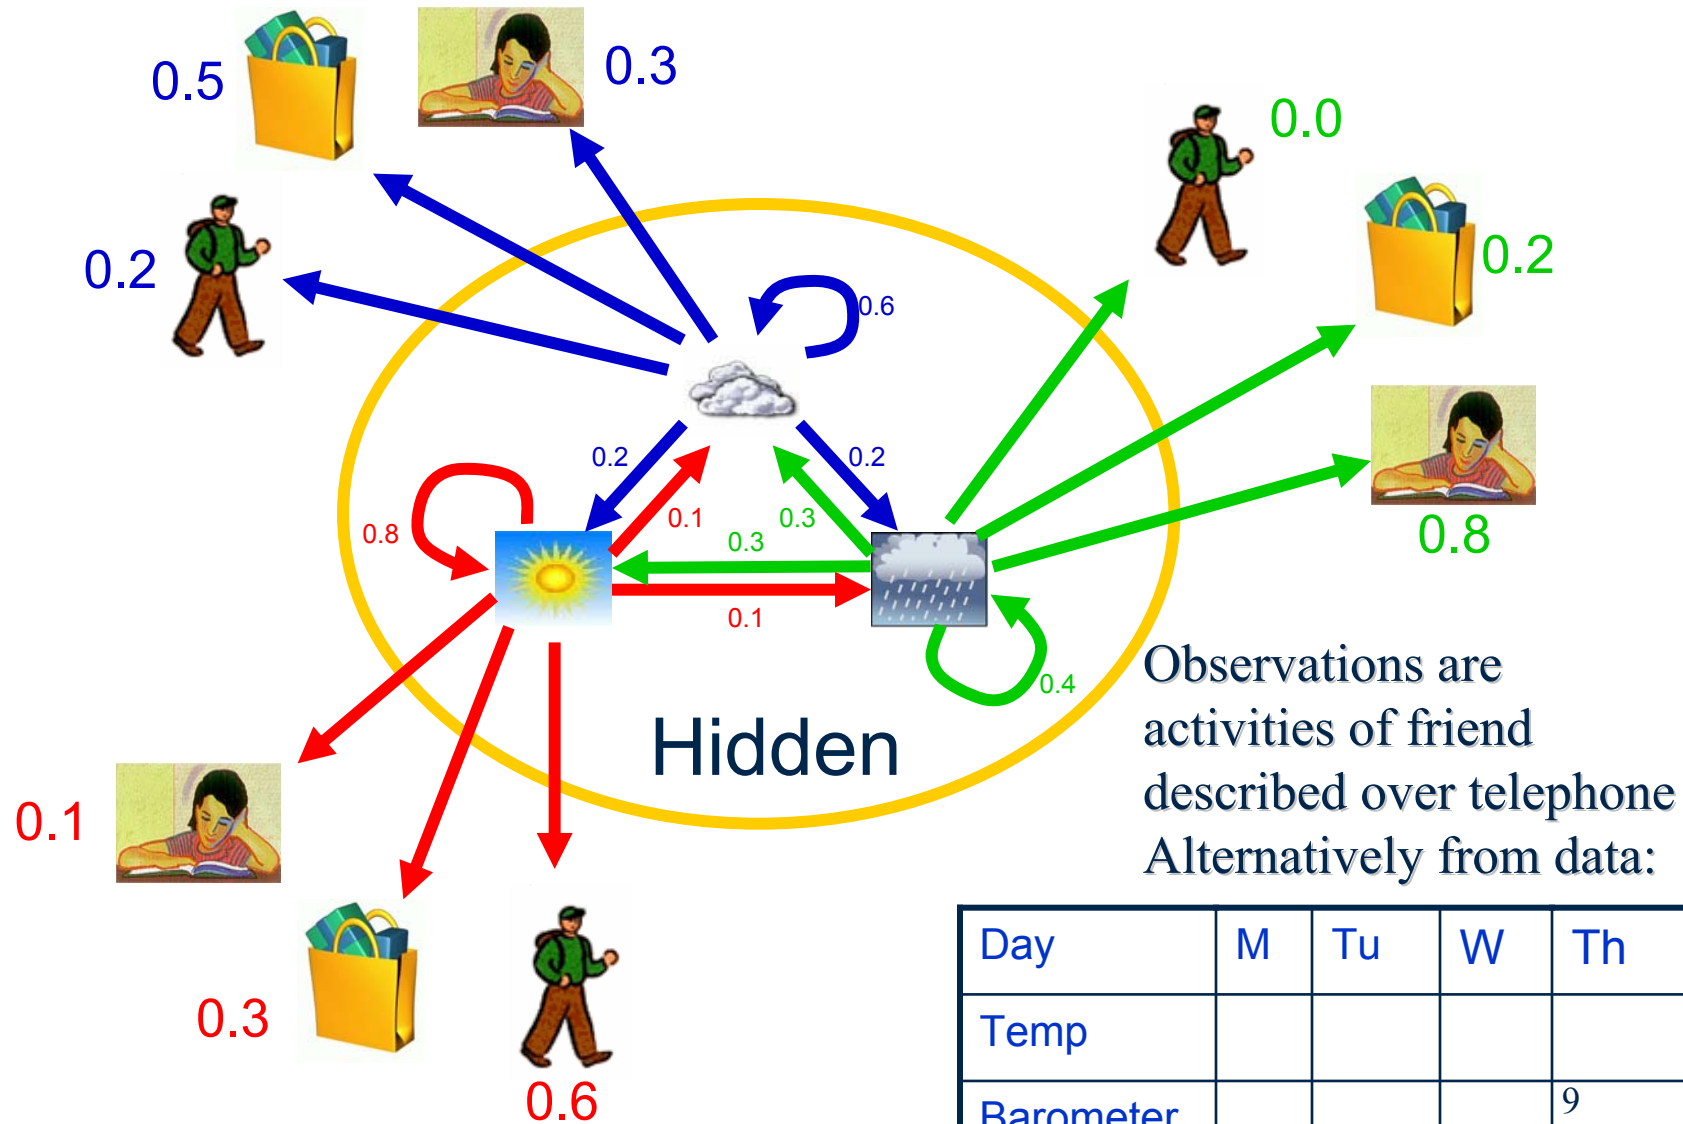
# Making a Sequence of Decisions

- Processes in time, states at time  $t$  are influenced by a state at time  $t-1$
- Wish to predict next value from previous values, e.g., financial forecasting
- Impractical to consider general dependence of future dependence on all previous observations
  - Complexity grows without limit as number of observations increases
- Markov models assume dependence on most recent observations

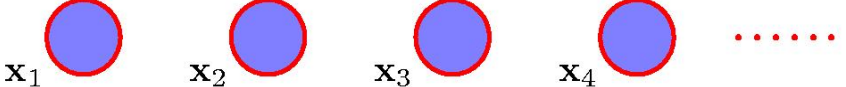
# Latent Variables

- While Markov models are tractable they are severely limited
- Introduction of latent variables provides a more general framework
- Lead to state-space models
- When latent variables are:
  - Discrete
    - they are called *Hidden Markov models*
  - Continuous
    - they are *linear dynamical systems*

# Hidden Markov Model



# Markov Model Assuming Independence

- Simplest model: 
  - Assume observations are independent
  - Graph without links
- To predict whether it rains tomorrow is only based on relative frequency of rainy days
- Ignores influence of whether it rained the previous day


# Markov Model

- Most general Markov model for observations  $\{x_n\}$
- Product rule to express joint distribution of sequence of observations

$$p(x_1, \dots, x_N) = \prod_{n=1}^N p(x_n \mid x_1, \dots, x_{n-1})$$

# First Order Markov Model

- Chain of observations  $\{x_n\}$



- Joint distribution for a sequence of  $n$  variables

$$p(x_1, \dots, x_N) = p(x_1) \prod_{n=2}^N p(x_n | x_{n-1})$$

- It can be verified (using product rule from above) that

$$p(x_n | x_1 \dots x_{n-1}) = p(x_n | x_{n-1})$$

- If model is used to predict next observation, distribution of prediction will only depend on preceding observation and independent of earlier observations
- Stationarity implies conditional distributions  $p(x_n | x_{n-1})$  are all equal

# Markov Model – Sequence probability

- What is the probability that the weather for the next 7 days will be “S-S-R-R-S-C-S”?

$$O = \{S_3, S_3, S_3, S_1, S_1, S_3, S_2, S_3\}$$


—Find the probability of O, given the model.

$$\begin{aligned} P(O \mid Model) &= P(S_3, S_3, S_3, S_1, S_1, S_3, S_2, S_3 \mid Model) \\ &= P(S_3) \cdot P(S_3 \mid S_3) \cdot P(S_3 \mid S_3) \cdot P(S_1 \mid S_3) \\ &\quad \cdot P(S_1 \mid S_1) \cdot P(S_3 \mid S_1) \cdot P(S_2 \mid S_3) \cdot P(S_3 \mid S_2) \\ &= \pi_3 \cdot a_{33} \cdot a_{33} \cdot a_{31} \cdot a_{11} \cdot a_{13} \cdot a_{32} \cdot a_{23} \\ &= 1 \cdot (0.8) \cdot (0.8) \cdot (0.1) \cdot (0.4) \cdot (0.3) \cdot (0.1) \cdot (0.2) \\ &= 1.536 \times 10^{-4} \end{aligned}$$



# Second Order Markov Model

- Conditional distribution of observation  $x_n$  depends on the values of two previous observations  $x_{n-1}$  and  $x_{n-2}$



$$p(x_1, \dots, x_N) = p(x_1) p(x_2 | x_1) \prod_{n=3}^N p(x_n | x_{n-1}, x_{n-2})$$


- Each observation is influenced by previous two observations

# $M^{th}$ Order Markov Source

- Conditional distribution for a particular variable depends on previous  $M$  variables
- Pay a price for number of parameters
- Discrete variable with  $K$  states
  - First order:  $p(x_n|x_{n-1})$  needs  $K-1$  parameters for each value of  $x_{n-1}$  for each of  $K$  states of  $x_n$  giving  $K(K-1)$  parameters
  - $M^{th}$  order will need  $K^{M-1}(K-1)$  parameters

# Introducing Latent Variables

- Model for sequences not limited by Markov assumption of any order but with limited number of parameters
- For each observation  $x_n$ , introduce a latent variable  $z_n$
- $z_n$  may be of different type or dimensionality to the observed variable
- Latent variables form the Markov chain
- Gives the “state-space model”




# Conditional Independence with Latent Variables

- Satisfies key assumption that


$$z_{n+1} \perp z_{n-1} \mid z_n$$

- From d-separation

When latent node  $z_n$  is filled, the only path between  $z_{n-1}$  and  $z_{n+1}$  has a head-to-tail node that is blocked



# Jt Distribution with Latent Variables




- Joint distribution for this model

$$p(x_1, \dots, x_N, z_1, \dots, z_n) = p(z_1) \left[ \prod_{n=2}^N p(z_n | z_{n-1}) \right] \prod_{n=1}^N p(x_n | z_n)$$

- There is always a path between any  $x_n$  and  $x_m$  via latent variables which is never blocked
- Thus predictive distribution  $p(x_{n+1} | x_1, \dots, x_n)$  for observation  $x_{n+1}$  does not exhibit conditional independence properties and is hence dependent on all previous observations

# Two Models Described by Graph



1. Hidden Markov Model: If latent variables are discrete:  
Observed variables in a HMM may be discrete or continuous
2. Linear Dynamical Systems: If both latent and observed variables are Gaussian

# Further Topics on Sequential Data

- Hidden Markov Models:

<http://www.cedar.buffalo.edu/~srihari/CSE574/Chap11/Ch11.2-HiddenMarkovModels.pdf>

- Extensions of HMMs:

<http://www.cedar.buffalo.edu/~srihari/CSE574/Chap11/Ch11.3-HMMExtensions.pdf>

- Linear Dynamical Systems:

<http://www.cedar.buffalo.edu/~srihari/CSE574/Chap11/Ch11.4-LinearDynamicalSystems.pdf>

- Conditional Random Fields:

<http://www.cedar.buffalo.edu/~srihari/CSE574/Chap11/Ch11.5-ConditionalRandomFields.pdf>